

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import sklearn
```

## importing Dataset

```
In [3]: dataset = pd.read_csv('50_Startups.csv')
```

```
In [5]: dataset.head(),dataset.tail()
```

```
Out[5]: (  R&D Spend  Administration  Marketing Spend    Profit
0  165349.20      136897.80      471784.10  192261.83
1  162597.70      151377.59      443898.53  191792.06
2  153441.51      101145.55      407934.54  191050.39
3  144372.41      118671.85      383199.62  182901.99
4  142107.34       91391.77      366168.42  166187.94,
   R&D Spend  Administration  Marketing Spend    Profit
45   1000.23      124153.04      1903.93  64926.08
46   1315.46      115816.21     297114.46  49490.75
47     0.00      135426.92         0.00  42559.73
48    542.05       51743.15         0.00  35673.41
49     0.00      116983.80      45173.06  14681.40)
```

```
In [7]: print('rows : ',dataset.shape[0])
print('columns : ',dataset.shape[1])
```

```
rows : 50
columns : 4
```

```
In [8]: #checking for repeated values in dataset
dataset.duplicated().sum()
```

```
Out[8]: 0
```

```
In [9]: #checking for null values in dataset
dataset.isnull().sum()
```

```
Out[9]: R&D Spend      0
Administration      0
Marketing Spend      0
Profit              0
dtype: int64
```

```
In [10]: #dataset schema
dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50 entries, 0 to 49
Data columns (total 4 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   R&D Spend           50 non-null    float64
```

```

1 Administration 50 non-null float64
2 Marketing Spend 50 non-null float64
3 Profit 50 non-null float64
dtypes: float64(4)

```

```

In [12]: #corr function
#corr() is used to find the pairwise correlation of all columns in the dataframe.
c = dataset.corr()

```

In [13]:

c

```

Out[13]:

```


	R&D Spend	Administration	Marketing Spend	Profit
R&D Spend	1.000000	0.241955	0.724248	0.972900
Administration	0.241955	1.000000	-0.032154	0.200717
Marketing Spend	0.724248	-0.032154	1.000000	0.747766
Profit	0.972900	0.200717	0.747766	1.000000

correlation matrix

```

In [14]: sns.heatmap(c,annot=True,cmap='Blues')
plt.show()

```



Outliers detection in the target variable


```

In [15]: outliers = ['Profit']
plt.rcParams['figure.figsize'] = [8,8]
sns.boxplot(data=dataset[outliers], orient="v", palette="Set2", width=0.7) # ori
# ori

plt.title("Outliers Variable Distribution")
plt.ylabel("Profit Range")
plt.xlabel("Continuous Variable")

plt.show()


```



```
In [18]: sns.distplot(dataset['Profit'], bins=5, kde=True)
plt.show()
```

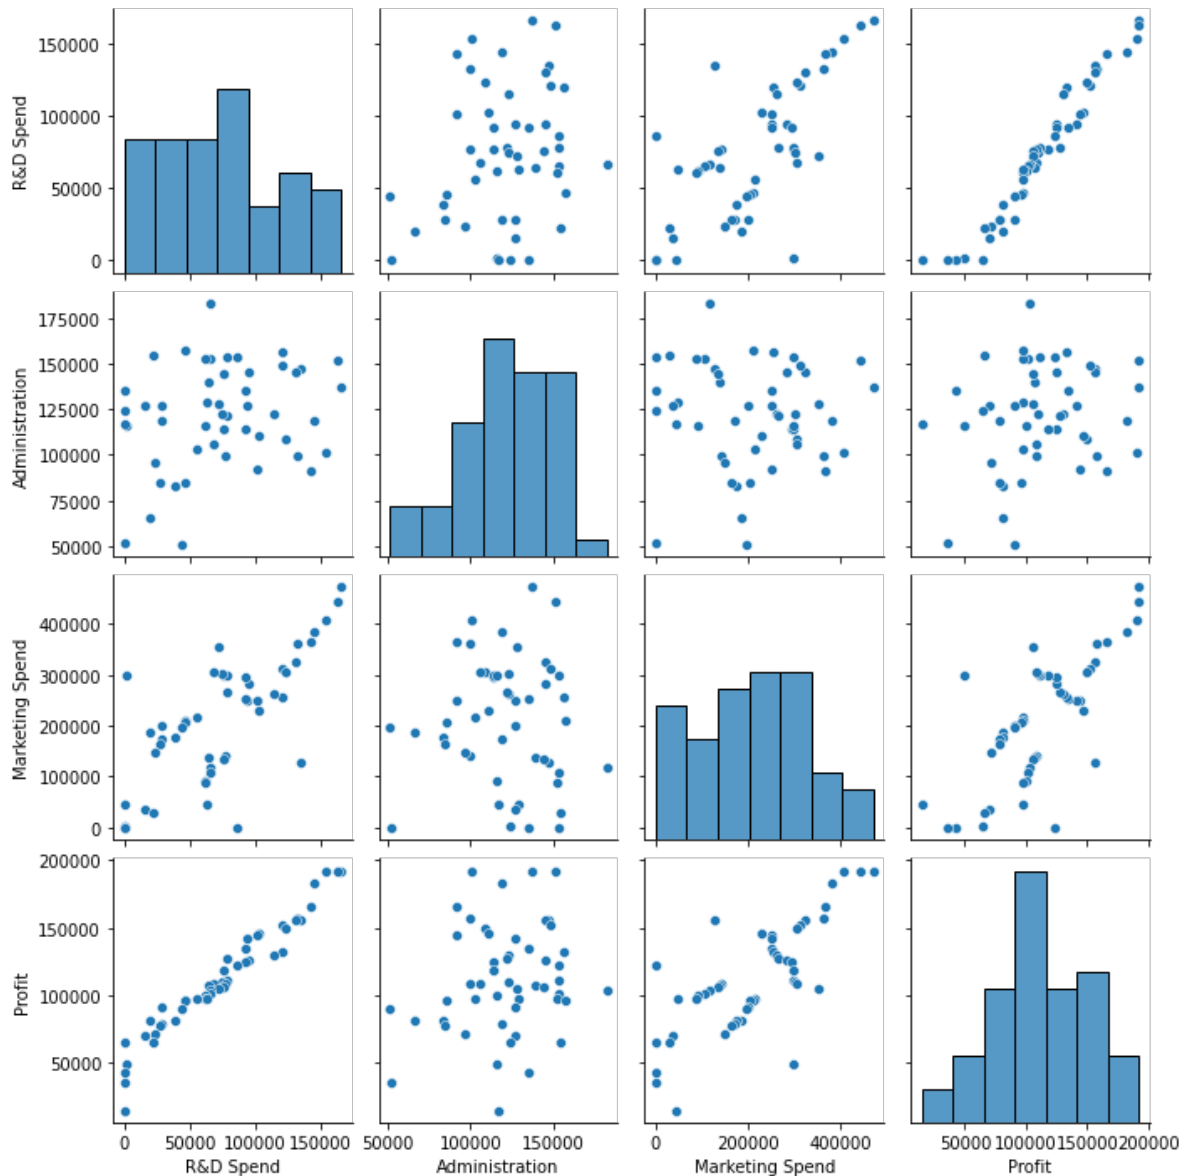
c:\users\91755\appdata\local\programs\python\python39\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```



In [19]:

```
sns.pairplot(dataset)
plt.figure(figsize=(5,5))
plt.show()
```



<Figure size 360x360 with 0 Axes>

## model preparation

In [21]:

```
# splitting Dataset in Dependent & Independent Variables
X = dataset.iloc[:, :-1].values
y = dataset.iloc[:, 3].values
```

```
In [24]: from sklearn.preprocessing import LabelEncoder
```

```
In [26]: #Label Encoder : Encode Labels with value between 0 and n_classes-1.
labelencoder = LabelEncoder()
X[:, 2] = labelencoder.fit_transform(X[:, 2])
X1 = pd.DataFrame(X)
X1.head()
```

```
Out[26]:
```

	0	1	2
0	165349.20	136897.80	47.0
1	162597.70	151377.59	46.0
2	153441.51	101145.55	45.0
3	144372.41	118671.85	44.0
4	142107.34	91391.77	43.0

## split the data into training and testing data

```
In [27]: split the data into training and testing data
from sklearn.model_selection import train_test_split

x_train,x_test,y_train,y_test = train_test_split(X,y,train_size=0.7,random_state=0)
x_train
```

```
Out[27]: array([[1.3029813e+05, 1.4553006e+05, 4.0000000e+01],
 [1.1994324e+05, 1.5654742e+05, 2.8000000e+01],
 [1.0002300e+03, 1.2415304e+05, 1.0000000e+00],
 [5.4205000e+02, 5.1743150e+04, 0.0000000e+00],
 [6.5605480e+04, 1.5303206e+05, 8.0000000e+00],
 [1.1452361e+05, 1.2261684e+05, 2.9000000e+01],
 [6.1994480e+04, 1.1564128e+05, 7.0000000e+00],
 [6.3408860e+04, 1.2921961e+05, 5.0000000e+00],
 [7.8013110e+04, 1.2159755e+05, 3.0000000e+01],
 [2.3640930e+04, 9.6189630e+04, 1.4000000e+01],
 [7.6253860e+04, 1.1386730e+05, 3.4000000e+01],
 [1.5505730e+04, 1.2738230e+05, 3.0000000e+00],
 [1.2054252e+05, 1.4871895e+05, 3.9000000e+01],
 [9.1992390e+04, 1.3549507e+05, 2.7000000e+01],
 [6.4664710e+04, 1.3955316e+05, 1.2000000e+01],
 [1.3187690e+05, 9.9814710e+04, 4.2000000e+01],
 [9.4657160e+04, 1.4507758e+05, 3.1000000e+01],
 [2.8754330e+04, 1.1854605e+05, 1.6000000e+01],
 [0.0000000e+00, 1.1698380e+05, 4.0000000e+00],
 [1.6259770e+05, 1.5137759e+05, 4.6000000e+01],
 [9.3863750e+04, 1.2732038e+05, 2.6000000e+01],
 [4.4069950e+04, 5.1283140e+04, 1.9000000e+01],
 [7.7044010e+04, 9.9281340e+04, 1.3000000e+01],
 [1.3461546e+05, 1.4719887e+05, 1.0000000e+01],
 [6.7532530e+04, 1.0575103e+05, 3.7000000e+01],
 [2.8663760e+04, 1.2705621e+05, 2.0000000e+01],
 [7.8389470e+04, 1.5377343e+05, 3.5000000e+01],
 [8.6419700e+04, 1.5351411e+05, 0.0000000e+00],
 [1.2333488e+05, 1.0867917e+05, 3.8000000e+01],
 [3.8558510e+04, 8.2982090e+04, 1.7000000e+01],
 [1.3154600e+03, 1.1581621e+05, 3.3000000e+01],
```

```
[1.4437241e+05, 1.1867185e+05, 4.4000000e+01],
[1.6534920e+05, 1.3689780e+05, 4.7000000e+01],
[0.0000000e+00, 1.3542692e+05, 0.0000000e+00],
.....
```

In [30]: *# Feature Scaling -- Useful when Features have different units*

```
"""from sklearn.preprocessing import StandardScaler
sc_X = StandardScaler()
X_train = sc_X.fit_transform(X_train)
X_test = sc_X.transform(X_test)
sc_y = StandardScaler()
y_train = sc_y.fit_transform(y_train)
y_test = sc_y.fit_transform(y_test)"""
```

Out[30]: 'from sklearn.preprocessing import StandardScaler\nsc\_X = StandardScaler()\nX\_train = sc\_X.fit\_transform(X\_train)\nX\_test = sc\_X.transform(X\_test)\nsc\_y = StandardScaler()\ny\_train = sc\_y.fit\_transform(y\_train)\ny\_test = sc\_y.fit\_transform(y\_test)'

In [31]: **from** sklearn.linear\_model **import** LinearRegression

```
model = LinearRegression()
model.fit(x_train,y_train)
print('Model has been trained successfully')
```

Model has been trained successfully

In [32]: `y_pred = model.predict(x_test)`  
y\_pred

Out[32]: array([103365.65430448, 132409.63159464, 133669.58924177, 71596.33493623, 179574.8809234 , 114195.96899299, 65656.85292429, 97938.81018901, 114412.29898539, 169772.36831918, 96050.9051499 , 87515.25731045, 110242.6075272 , 90000.89195708, 127479.23515393])

In [33]: `testing_data_model_score = model.score(x_test, y_test)`  
`print("Model Score/Performance on Testing data",testing_data_model_score)`  
  
`training_data_model_score = model.score(x_train, y_train)`  
`print("Model Score/Performance on Training data",training_data_model_score)`

Model Score/Performance on Testing data 0.9324057207634493  
Model Score/Performance on Training data 0.9506671824404848

In [34]: `df = pd.DataFrame(data={'Predicted value':y_pred.flatten(),'Actual Value':y_test.flatten()})`  
df

Out[34]:

	Predicted value	Actual Value
--	-----------------	--------------

0	103365.654304	103282.38
1	132409.631595	144259.40
2	133669.589242	146121.95
3	71596.334936	77798.83
4	179574.880923	191050.39
5	114195.968993	105008.31

	Predicted value	Actual Value
6	65656.852924	81229.06
7	97938.810189	97483.56
8	114412.298985	110352.25
9	169772.368319	166187.94
10	96050.905150	96778.92
11	87515.257310	96479.51
12	110242.607527	105733.54
13	90000.891957	96712.80

## regression metrics

R2 score

```
In [35]: from sklearn.metrics import r2_score
r2Score = r2_score(y_pred, y_test)
print("R2 score of model is : ", r2Score*100)
```

R2 score of model is : 93.21346390789374

mean squared error

```
In [36]: from sklearn.metrics import mean_squared_error
mse = mean_squared_error(y_pred, y_test)
print("Mean Squarred Error is : ", mse*100)
```

Mean Squarred Error is : 6524519362.317416

root Mean squared Error

```
In [37]: rmse = np.sqrt(mean_squared_error(y_pred, y_test))
print("Root Mean Squarred Error is : ", rmse*100)
```

Root Mean Squarred Error is : 807744.9697966195

Mean absolute Error

```
In [38]: from sklearn.metrics import mean_absolute_error

mae = mean_absolute_error(y_pred, y_test)
print("Mean Absolute Error is : ", mae)
```

Mean Absolute Error is : 6603.238628961085

In [ ]: